

DATA SCIENCE & MACHINE LEARNING

Introduction Course

(23 JAN - 01 March 2019)

Coordination Team:

Augusto Albuquerque (ISCTE-IUL)
Inês Cortez Esteves (Deloitte)
Pedro Sebastião (ISCTE-IUL/AUDAX-ISCTE)
Pedro Tavares (Deloitte)

Team:

António Raimundo (ISCTE-IUL)
Dária Baikova (ISCTE-IUL)
João Oliveira (ISCTE-IUL)
Ricardo Ribeiro (ISCTE-IUL)
Sérgio Moro (ISCTE-IUL)

DATA SCIENCE & MACHINE LEARNING

Introduction Course

General Contents (23 JAN 2019)

- **#0 Introduction – Overview (23 JAN 2019)**
- **#1 IPython – Beyond Normal Python (23 JAN 2019)**
- #2 Introduction to NumPy
- #3 Data Manipulation and Visualization
- #4 Statistical Analysis
- #5 Time-Series and High Performance Pandas
- #6 Introduction to Machine Learning
- #7 Machine Learning Techniques
- #8 Deep Learning
- #9 Network Analysis Social Networks
- #10 Recommendation Systems using NLP

DATA SCIENCE & MACHINE LEARNING

#0 Introduction – Overview

(23 JAN 2019)


Lecturers: Sérgio Moro & António Raimundo

#0 Introduction – Overview

Contents (23 JAN 2019)

- Introduction to Data Science
- CRISP-DM
- Practical example: CRM
- Practical example: Service Desk System

#0 Introduction – Overview




What is Data Science?

Interdisciplinary data-driven approach to address problems for which there is **available data**.

#0 Introduction – Overview

Data Mining CRISP-DM Methodology



Data Mining: knowledge **discovery** process that aims to unveil **insightful patterns** from **raw data** (it is encompassed within Data Science)

Text Mining: Unstructured data

Note: Machine Learning (ML) aims to learn from data.

Thus, ML may (and it is often) used for building models within a DM project.

#0 Introduction – Overview

Data Mining CRISP-DM Methodology

Business Understanding

- Hear from business experts:
 - How do they define the problem?
 - What are the factors that may influence the result?
 - How would they measure the success of a possible solution?
- Create a project
 - Goals
 - Risks
 - Stakeholders
 - Plan

#0 Introduction – Overview

Data Mining CRISP-DM Methodology

Data Understanding

- What are the variables that best characterize the problem (input variables)? What is their range of possible values and type (e.g., categorical vs. numeric)?
- From these, which best translate the business goal (for supervised learning)?
- Are there data quality issues?
 - Inconsistent data originated in more than one source
 - Missing values

#0 Introduction – Overview

Data Mining CRISP-DM Methodology

Data Preparation

- Data integration: consolidate data from distinct sources
- Feature selection: which input variables?
- Feature Engineering: computing new interesting variables
- Data cleaning: deal with missing values, outliers

#0 Introduction – Overview

Data Mining CRISP-DM Methodology

Modeling

- Uses Statistics or Machine Learning algorithms to train a model
- It depends on the goal:
 - Unsupervised learning – find relations between input variables
 - Supervised learning – use a target (dependent) variable to build a model based on input variables
 - If the target is categorical (e.g., grant loan: Y/N), then it is a classification problem; else (for numeric variables, e.g., how many products sold), it is a regression problem

#0 Introduction – Overview

Data Mining CRISP-DM Methodology

Evaluation

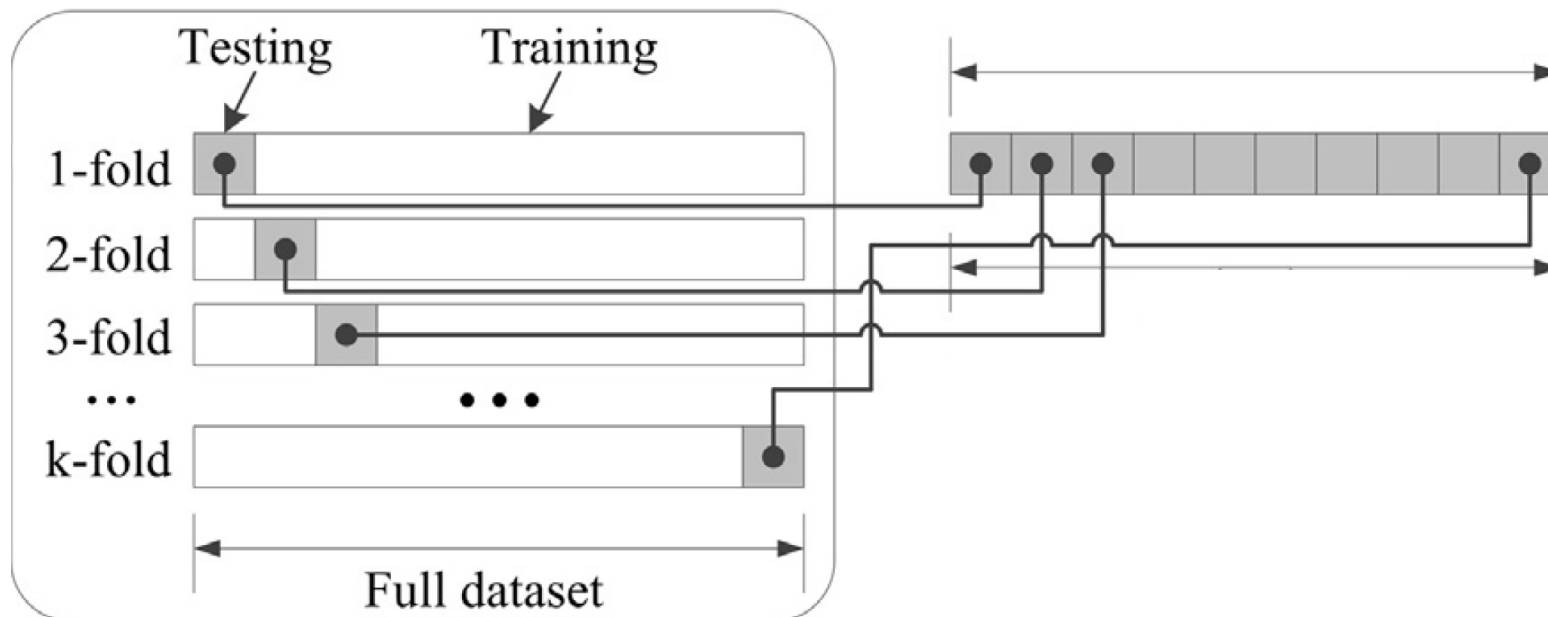
- Assess if the model is accurate
- For supervised learning models: how close are the predictions from the real values?
- For unsupervised learning: how consistent are the identified groups?
- In a production scenario: a model is trained using past data, and it is evaluated using future data
- During project (or for a quick model tuning): partition the dataset into train (to build the model)and test sets

#0 Introduction – Overview

Data Mining CRISP-DM Methodology

Evaluation

- K-fold cross validation: each fold is used once for testing and K-1 times for training



#0 Introduction – Overview

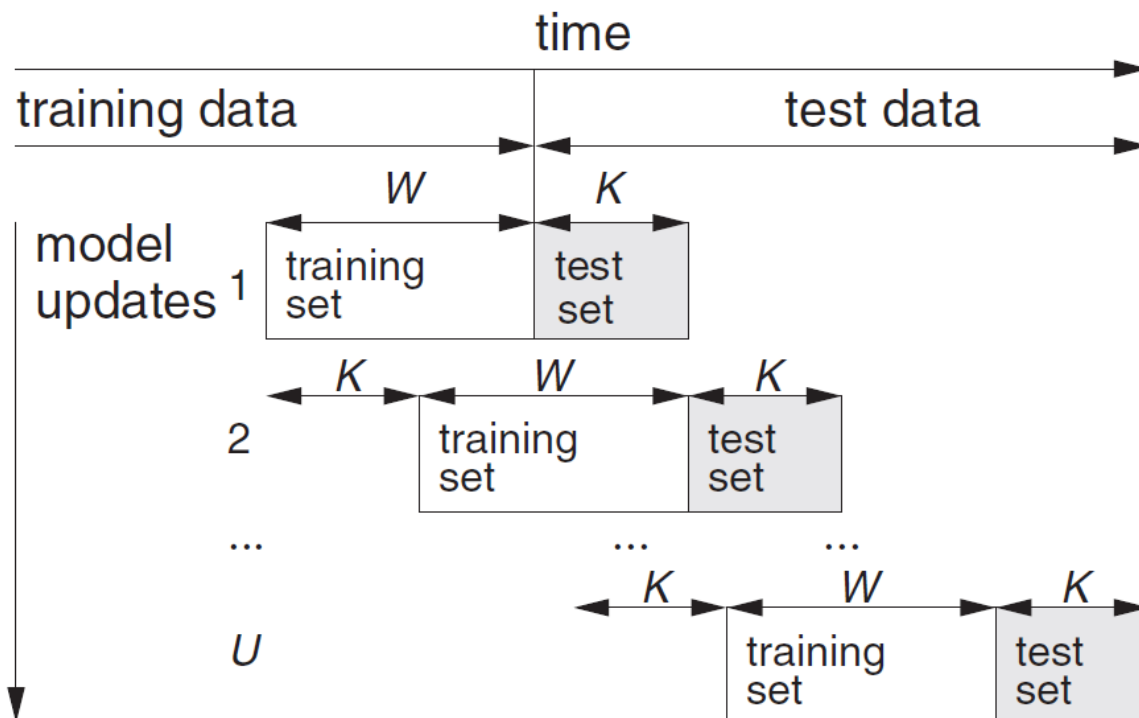
Data Mining CRISP-DM Methodology

Evaluation

- Rolling Windows: to simulate a periodic train/test scenario

All occurrences of the problem (also called instances) must be sorted by date

This is what will happen once the model is deployed into production




#0 Introduction – Overview

Data Mining CRISP-DM Methodology

Deployment

- Install the model into a production environment
- Yet, the work isn't over:
 - Real problems are dynamic and the model's performance can quickly deteriorate
 - New variables may need to be incorporated, or some removed
 - Out-of-the-ordinary events may induce unexpected behavior
- **Thus, tuning should be done periodically**
- CRISP-DM is a cyclic methodology that advocates revision iterations




#0 Introduction – Overview

Data Mining CRISP-DM Methodology

Deployment

- Install the model into a production environment
- Yet, the work isn't over:
 - Real problems are dynamic and the model's performance can quickly deteriorate
 - New variables may need to be incorporated, or some removed
 - Out-of-the-ordinary events may induce unexpected behavior
- **Thus, tuning should be done periodically**
- CRISP-DM is a cyclic methodology that advocates revision iterations




#0 Introduction – Overview

Customer Relationship Management (CRM)

Back Office
Analytical CRM

Usual questions:

- Which customers to contact?
- Which products to sell?



Front Office
Operational CRM

Usual operations:

- Contact customer
- Sell product



Customer


#0 Introduction – Overview

Customer Relationship Management (CRM)

Back Office
Analytical CRM

How to answer usual questions:

- Through implemented Data Mining models



Front Office
Operational CRM

How to make usual operations:

- Through operational applications (e.g., contact center)



Customer

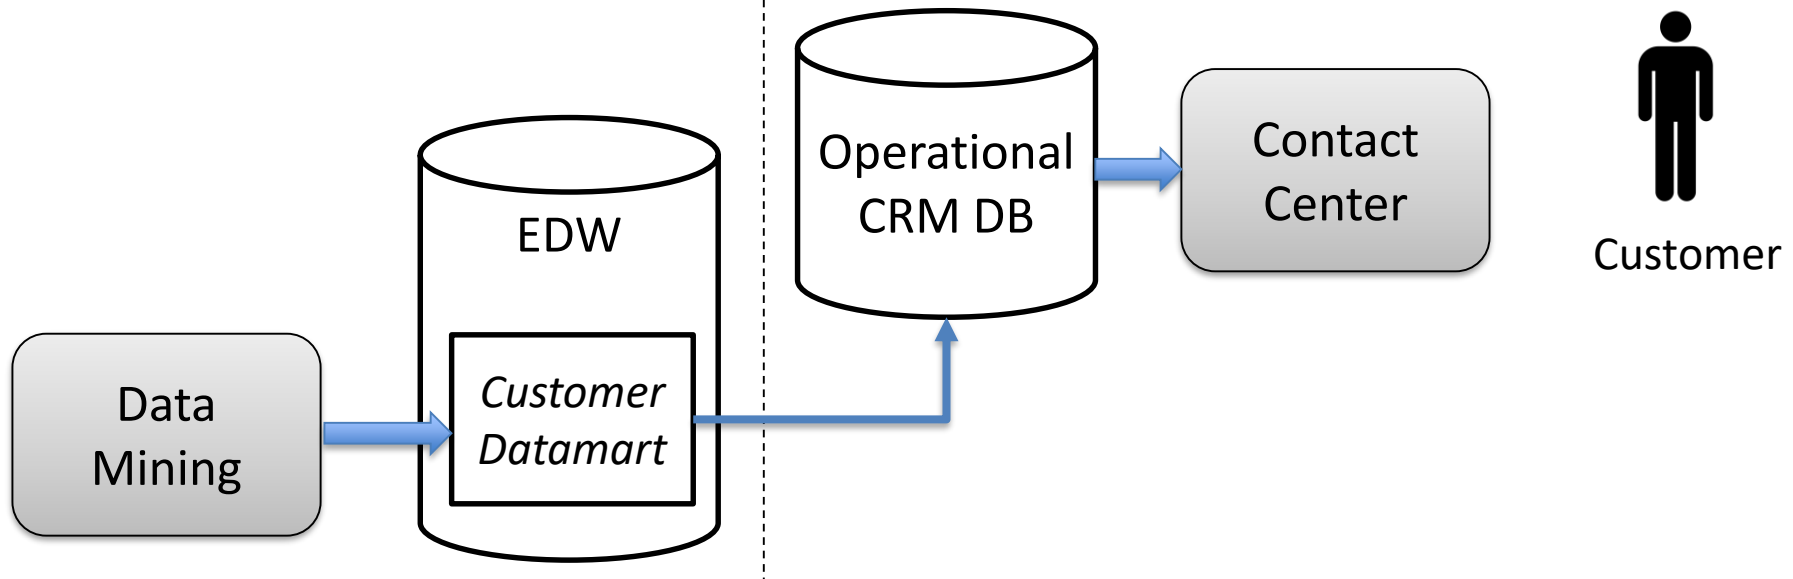
#0 Introduction – Overview

Customer Relationship Management (CRM)

Which customers to contact for selling a product?

Back Office
Analytical CRM

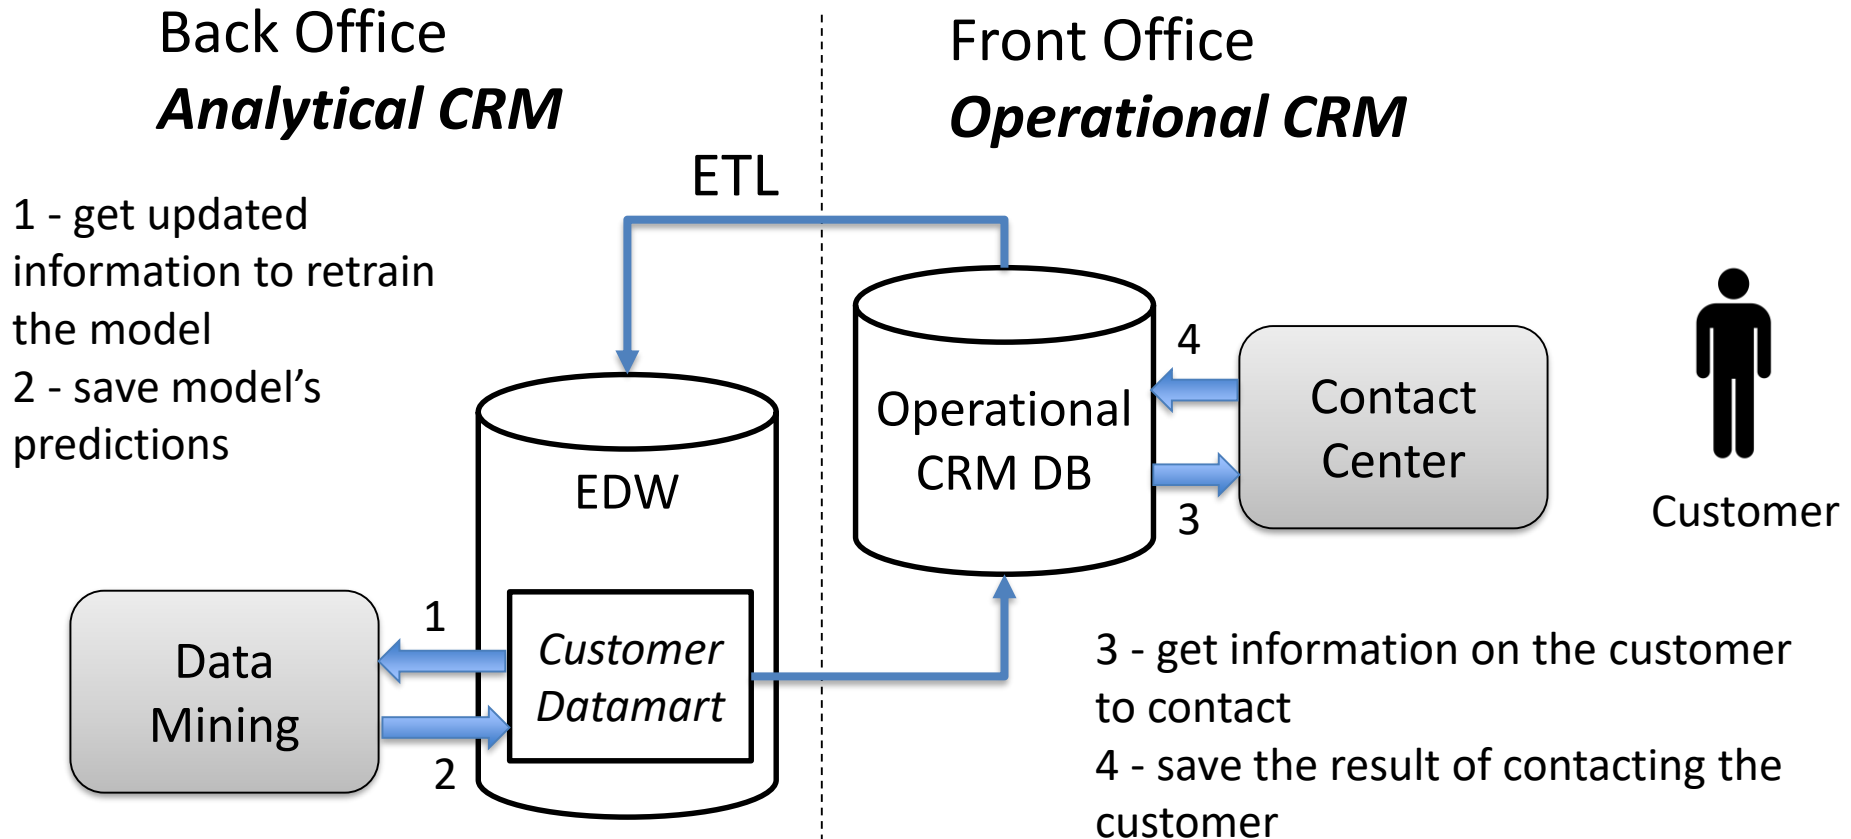
Front Office
Operational CRM



#0 Introduction – Overview

Customer Relationship Management (CRM)

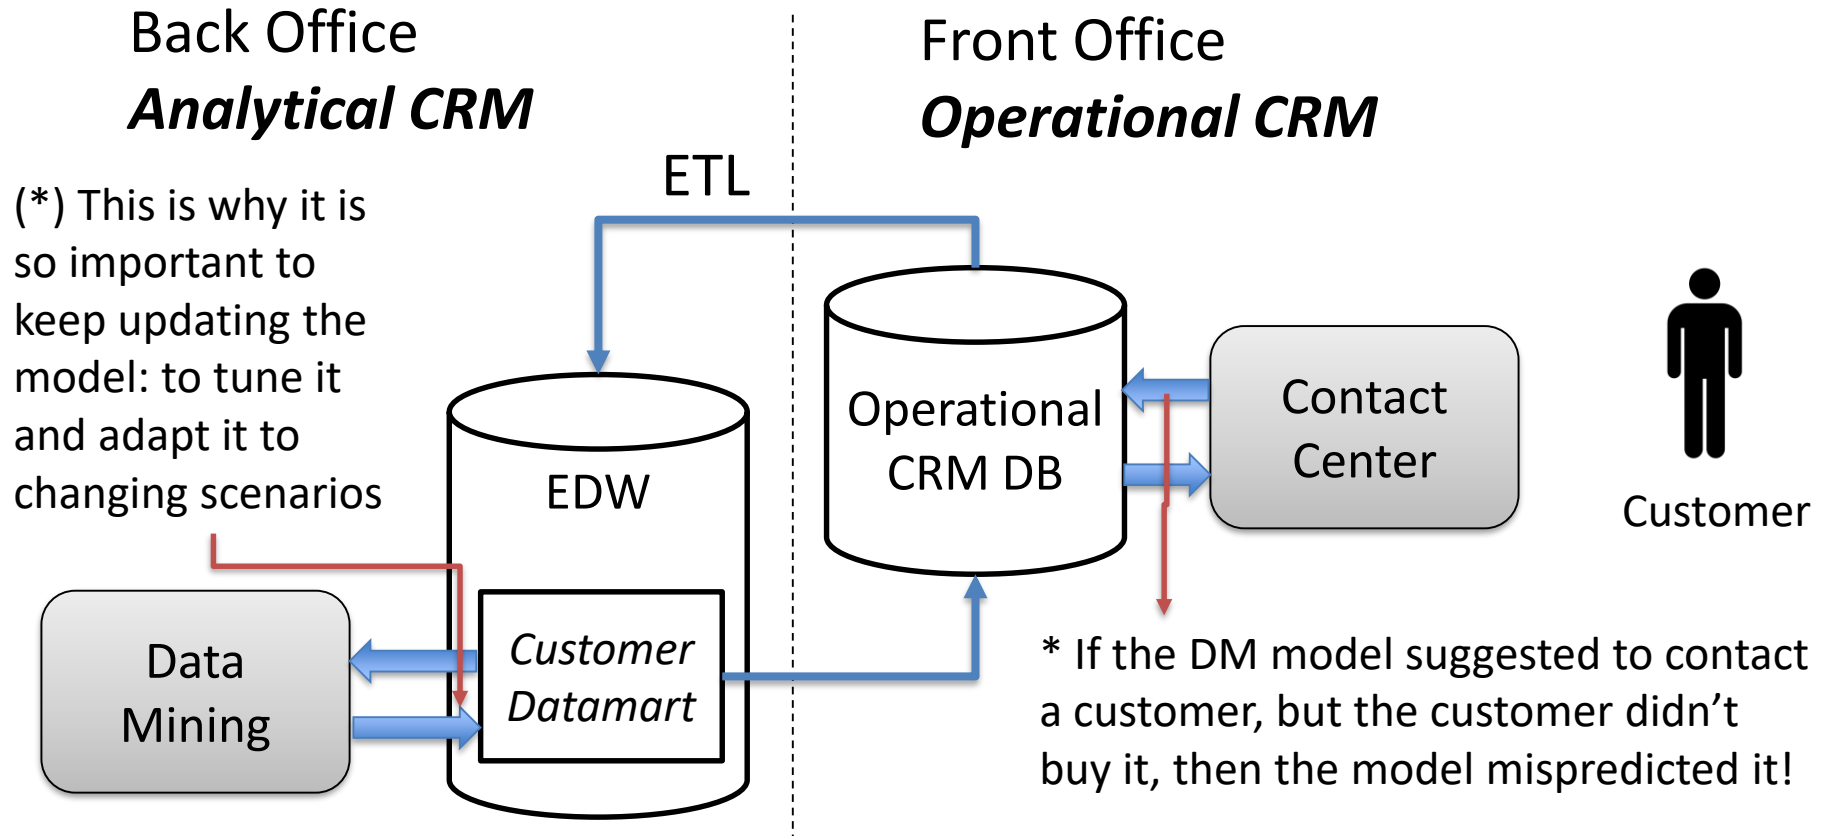
Closing the loop...



#0 Introduction – Overview

Customer Relationship Management (CRM)

Which customers to contact for selling a product?



#0 Introduction – Overview


Service Desk System (SD)

Incident management

Back Office

Usual questions:

- How to categorize incoming incidents (to prioritize them)?
- How to diagnose to suggest a solution?



Front Office

Usual operations:

- Categorize (or revise priority of) incidents
- Propose a solution to the problem
- Escalate to specialized staff

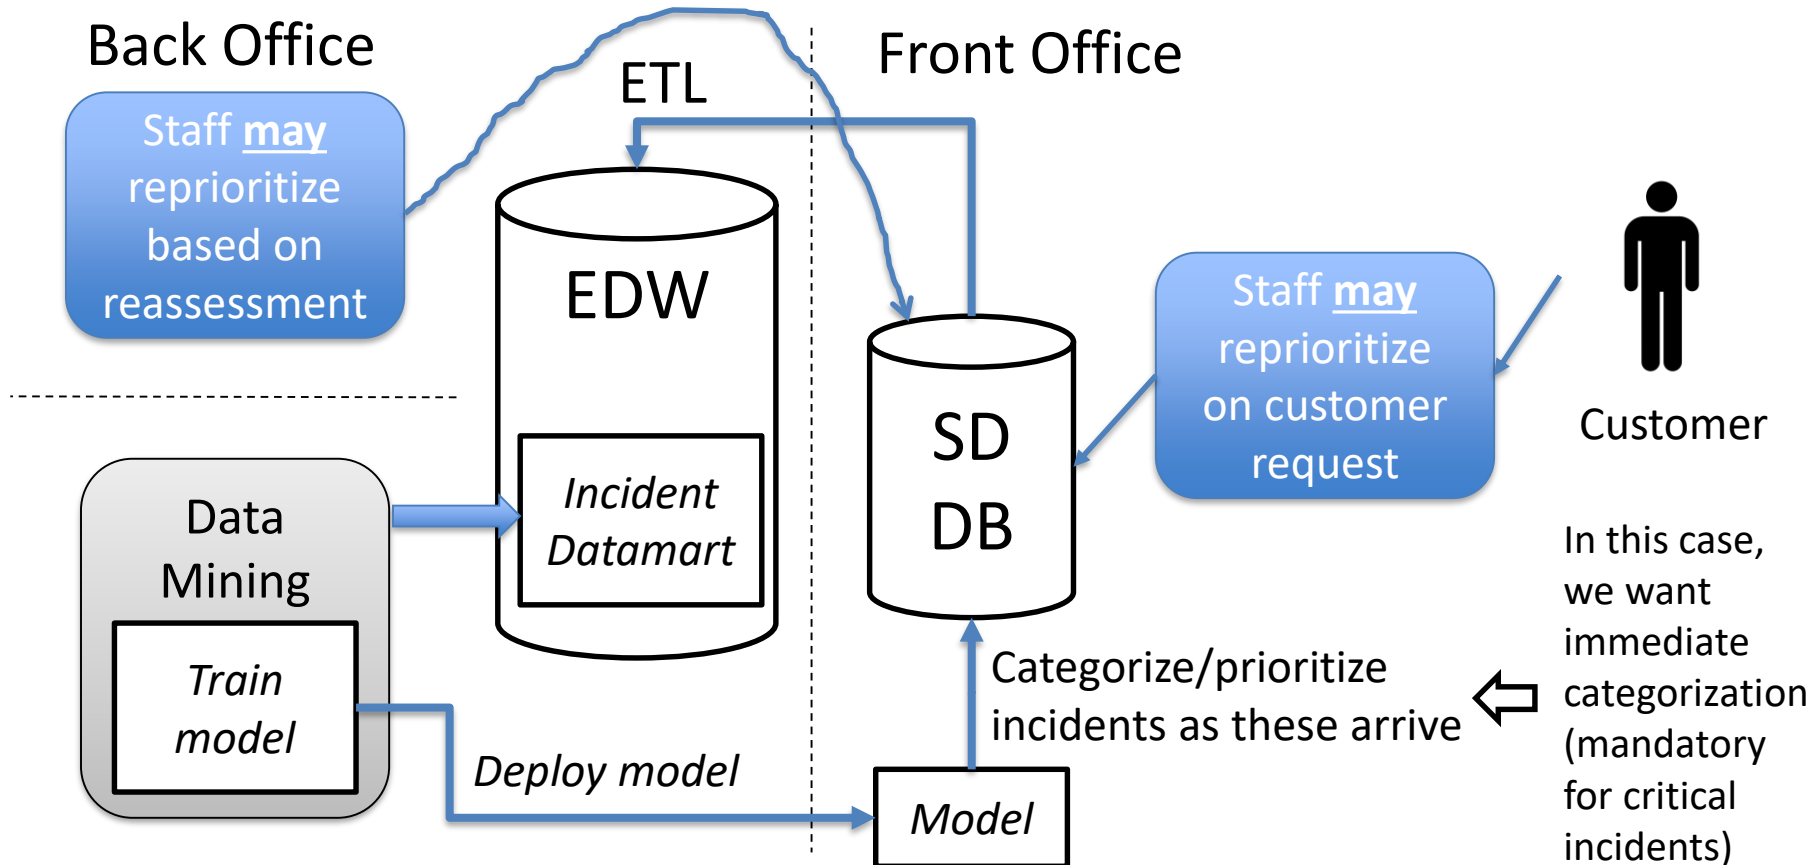


Customer

#0 Introduction – Overview

Service Desk System (SD)

Incident management – Automatic categorization



THANK YOU FOR YOUR PARTICIPATION

#0 Introduction – Overview

(23 JAN 2019)

Next module is:
#1 IPython – Beyond Normal Python
(23 JAN 2019)

DATA SCIENCE & MACHINE LEARNING

#1 IPython – Beyond Normal Python

(23 JAN 2019)

Lecturers: Sérgio Moro & António Raimundo

#1 IPython – Beyond Normal Python

Contents (23 JAN 2019)

- Shell or Notebook
- Help and Documentation in IPython
- Magic Commands
- Input and Output History
- Notebook and Shell Commands
- Errors and Debugging
- Profiling and Timing Code

#1 IPython – Beyond Normal Python

- **What is IPython?**
 - Features
- **Ties with Jupyter project**
 - Notebook
- **Shell or Notebook?**

```
In [9]: display(i)
```

IP[y]: IPython
Interactive Computing

```
In [3]: from IPython.display import SVG  
SVG(filename='python-logo.svg')
```

```
Out[3]:
```



We will work all the code examples on Jupyter Notebook!

#1 IPython – Beyond Normal Python

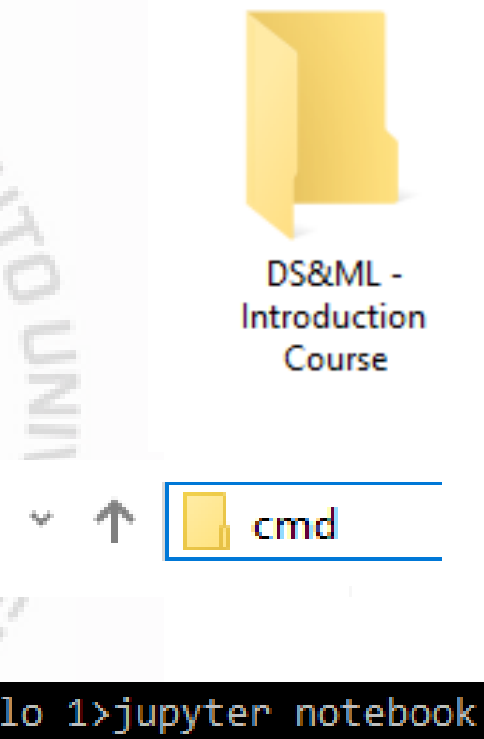
Jupyter Notebook



#1 IPython – Beyond Normal Python

• Launch Jupyter Notebook

- Create a folder on a desired location on your PC named: **“DS&ML – Introduction Course”**
- Create a subfolder called **“Module 1”**
 - You can create subfolders later for each module, according to the ongoing lectured module (2,3,4 and so on)
- (*Windows users*)
 - Enter **“Module 1”** folder, and type **“cmd”** on ‘*current location*’ bar
 - On the command line terminal, type: **‘jupyter notebook’**
- (*Linux/MacOS users*)
 - Open Terminal and type: **‘jupyter notebook’**
 - Navigate to **“Module 1”** folder using Notebook’s explorer interface.



#1 IPython – Beyond Normal Python

- **#1.1 Help and Documentation in IPython**
 - Accessing Documentation with ‘?’
 - Accessing Source Code with ‘??’
 - Exploring Modules with **Tab-Completion**
 - Tab-completion of object contents
 - Tab completion when importing

#1 IPython – Beyond Normal Python

- #1.2 Keyboard Shortcuts
 - Interactive Demonstration

Command Mode (press `Esc` to enable)

Edit Shortcuts

Edit Mode (press `Enter` to enable)

`F`: find and replace
`Ctrl-Shift-F`: open the command palette
`Ctrl-Shift-P`: open the command palette
`Enter`: enter edit mode
`P`: open the command palette
`Shift-Enter`: run cell, select below
`Ctrl-Enter`: run selected cells
`Alt-Enter`: run cell and insert below
`Y`: change cell to code
`M`: change cell to markdown
`R`: change cell to raw
`1`: change cell to heading 1
`2`: change cell to heading 2
`3`: change cell to heading 3
`4`: change cell to heading 4
`5`: change cell to heading 5
`6`: change cell to heading 6
`K`: select cell above
`Up`: select cell above
`Down`: select cell below
`J`: select cell below
`Shift-K`: extend selected cells above
`Shift-Up`: extend selected cells above

`Shift-Down`: extend selected cells below
`Shift-J`: extend selected cells below
`A`: insert cell above
`B`: insert cell below
`X`: cut selected cells
`C`: copy selected cells
`Shift-V`: paste cells above
`V`: paste cells below
`Z`: undo cell deletion
`D, D`: delete selected cells
`Shift-M`: merge selected cells, or current cell with cell below if only one cell is selected
`Ctrl-S`: Save and Checkpoint
`S`: Save and Checkpoint
`L`: toggle line numbers
`O`: toggle output of selected cells
`Shift-O`: toggle output scrolling of selected cells
`H`: show keyboard shortcuts
`I, I`: interrupt the kernel
`0, 0`: restart the kernel (with dialog)
`Esc`: close the pager
`Q`: close the pager
`Shift-L`: toggles line numbers in all cells,

`Tab`: code completion or indent
`Shift-Tab`: tooltip
`Ctrl-J`: indent
`Ctrl-I`: dedent
`Ctrl-A`: select all
`Ctrl-Z`: undo
`Ctrl-/`: comment
`Ctrl-D`: delete whole line
`Ctrl-U`: undo selection
`Insert`: toggle overwrite flag
`Ctrl-Home`: go to cell start
`Ctrl-Up`: go to cell start
`Ctrl-End`: go to cell end
`Ctrl-Down`: go to cell end
`Ctrl-Left`: go one word left

`Ctrl-Right`: go one word right
`Ctrl-Backspace`: delete word before
`Ctrl-Delete`: delete word after
`Ctrl-Y`: redo
`Alt-U`: redo selection
`Ctrl-M`: enter command mode
`Ctrl-Shift-F`: open the command palette
`Ctrl-Shift-P`: open the command palette
`Esc`: enter command mode
`Shift-Enter`: run cell, select below
`Ctrl-Enter`: run selected cells
`Alt-Enter`: run cell and insert below
`Ctrl-Shift-Minus`: split cell at cursor
`Ctrl-S`: Save and Checkpoint
`Down`: move cursor down
`Up`: move cursor up

#1 IPython – Beyond Normal Python

- **#1.3 Notebook's Magic Commands**

- Run external code: **%run**
 - Create a file named '**script1.py**' with the following content
(continue on Notebook)

```
def square(x):  
    """square a number"""  
    return x ** 2  
for N in range(1, 4):  
    print(N, "squared is", square(N))
```

- General description of magic commands: **%magic**
- List all available magic commands: **%lsmagic**

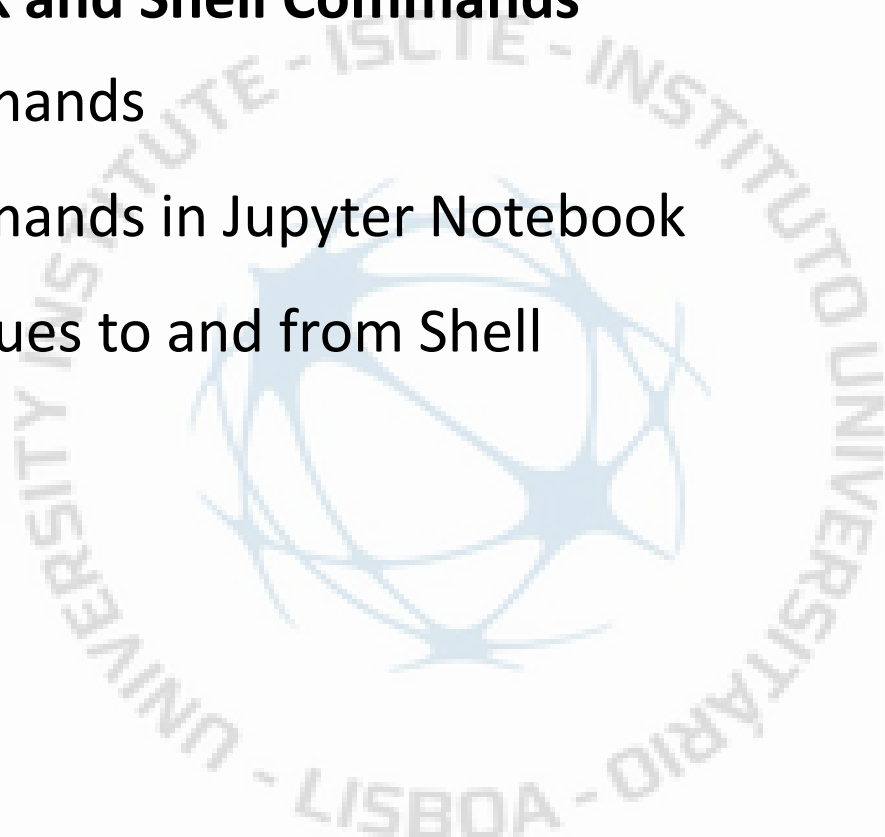
#1 IPython – Beyond Normal Python

- **#1.4 Input and Output History**

- Notebook's *In* and *Out* objects
- Underscore Shortcuts and Previous Outputs
- Suppressing Output
- Related Magic Commands

#1 IPython – Beyond Normal Python

- **#1.5 Notebook and Shell Commands**
 - Shell Commands
 - Shell Commands in Jupyter Notebook
 - Passing values to and from Shell



#1 IPython – Beyond Normal Python

• #1.6 Errors and Debugging

- Controlling Exceptions: **%xmode**
- Debug mode
- Partial list of debugging commands

Command	Description
list	Show the current location in the file
h(elp)	Show a list of commands, or find help on a specific command
q(uit)	Quit the debugger and the program
c(ontinue)	Quit the debugger, continue in the program
n(ext)	Go to the next step of the program
<enter>	Repeat the previous command
p(rint)	Print variables
s(tep)	Step into a subroutine
r(eturn)	Return out of a subroutine

#1 IPython – Beyond Normal Python

- **#1.7 Profiling and Timing Code**

- Timing Code Snippets: **%timeit** and **%time**
- Profiling Full Scripts: **%prun**
- Line-by-Line Profiling: **%lprun**
 - Install Line Profiler using '*pip*':
'pip install line_profiler'
- Profiling Memory Use: **%memit** and **%mprun**
 - Install Memory Profiler using '*pip*':
'pip install memory_profiler'

THANK YOU FOR YOUR PARTICIPATION

#1 IPython – Beyond Normal Python

(23 JAN 2019)

Next module is:
#2 Introduction to NumPy
(25 JAN 2019)